

# STAT 447 Assignment 1

Caden Hewlett

2024-01-11

## Question 1 : Sampling from a Joint Distribution

### Part 1 (INCOMPLETE)

Compute  $\mathbb{E}[(1 + Y_1)^X]$  mathematically (with a precise mathematical derivation).

**Solution:** We wish to compute  $\mathbb{E}(g, X, Y_1, \dots, Y_4)$ . In this case,  $g(x, y_1, \dots, y_4) = (1 + y_1)^x$ . We can hence write the expectation as follows:

$$\mathbb{E}[g(X, Y_1, \dots, Y_4)] = \sum_x \sum_{y_1} \sum_{y_2} \sum_{y_3} \sum_{y_4} g(x, y_1, \dots, y_4) p(x, y_1, y_2, y_3, y_4).$$

This can be simplified slightly as follows:

$$\mathbb{E}[(1 + Y_1)^X] = \sum_x \sum_{y_1} (1 + y_1)^x p(x, y_1, y_2, y_3, y_4)$$

Further, we know the specific  $y_2, y_3, y_4$  coin flip outcomes don't matter in this particular calculation, since our expectation is only in terms of  $X$  and  $Y_1$ . With this in mind, we can consider the following tree describing this situation.

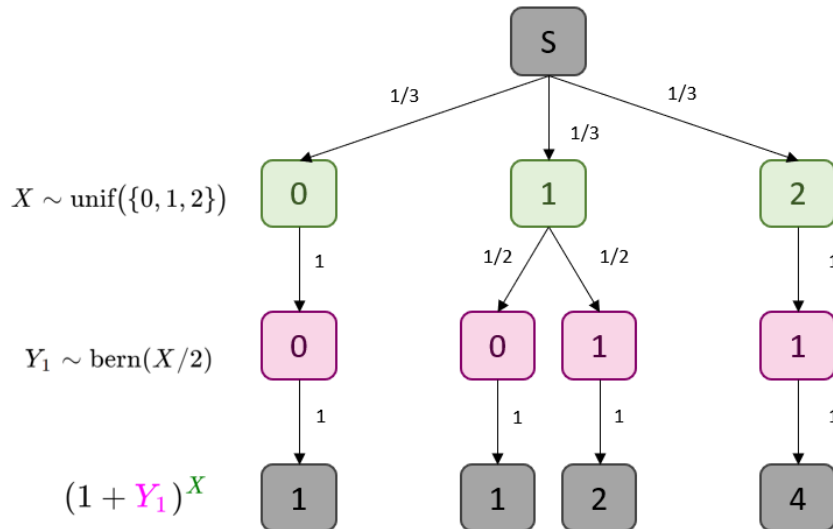


Figure 1: Decision Tree for First Coin Flip, with Function  $g(x, y_1)$

Then, applying the Law of the Unconscious Statistician (as seen in the bottom level of this tree), we can find  $\mathbb{E}((1 + Y_1)^X)$ . We use chain rule to propagate towards the probabilities aligning with the bottom leaves.

$$\mathbb{E}[(1 + Y_1)^X] = \frac{1}{3}(1) + \frac{1}{6}(1) + \frac{1}{6}(2) + \frac{1}{3}(4) = \boxed{\frac{13}{6}}$$

**Part 2** Write an R function called `forward_sample` that samples (“simulates”) from the joint distribution of  $(X, Y_1, Y_2, Y_3, Y_4)$ . As a general practice, fix the seed, and submit both the code and the output (here, a single sample).

```
library(extraDistr)
set.seed(19690720)
forward_sample <- function(){
  X = rdunif(1, 0, 2)
  return(
    c(X, sapply(1:4, function(x) rbinom(n = 1, size = 1, prob = X/2)))
  )
}
forward_sample()
```

```
## [1] 2 1 1 1 1
```

### Part 3

How can your code and the law of large number be used to approximate  $\mathbb{E}[(1 + Y_1)^X]$

Since the Law of Large numbers basically states that average of many *iid* random samples will approach the truth (if it exists), we can use our newly-created `forward_sample()` function to approximate via Monte Carlo simulation. We can conduct  $m$  simulations, and let  $\tilde{G}_M = (1 + y_1)^x$ . Then, the average given by  $\hat{G}_M = \frac{1}{M} \sum_{m=1}^M G^{(m)}$  is approximately equal to  $\mathbb{E}((1 + Y_1)^X)$ . The actual code to do this is below:

```
set.seed(19690720)
# declare number of iterations
num_iterations = 20000
# generate all simulations
simulations = sapply(1:20000, function(x){
  # generate a sample using the forward_sample() function
  sample_i = forward_sample()
  # calculate g(x, y) using the generated sample
  g_x = (1 + sample_i[2])^sample_i[1]
  # return this iteration
  return(g_x)
})
```

### Part 4

Compare the approximation from your code with you answer in part 1.

Here, we actually conduct the simulations and compare the results to those found by arithmetic operations.

```
# generate the estimate of E(g(X, Y))
simulated_mean = mean(simulations)
# value from simulations
print(paste("Simulated Value:", simulated_mean))
```

```
## [1] "Simulated Value: 2.1529"
```

```
# value from part 1
print(paste("Calculated Value:", round(13 / 6, 4)))
```

```
## [1] "Calculated Value: 2.1667"
```

```
# (observed - expected)/(expected) as a percent
print(paste("Percent Difference: ",
            round(abs(
                simulated_mean - 13 / 6
            ) / (13 / 6) * 100, 4), "%", sep = ' '))
```

```
## [1] "Percent Difference: 0.6354%"
```

So we can see that the result found from simulation is very close to that found in **Part 1**.

## Question 2 : Computing a Conditional

Suppose now that you observe the outcome of the 4 coin flips, but not the type of coin that was picked. Say you observe: “heads”, “heads”, “heads”, “heads” = [0, 0, 0, 0].

**Note:** We will use precise and careful notation to solve, not skipping any steps.

### Part 1

Write mathematically: “Given you observe 4 heads, what is the probability that you picked the standard coin?”

We know that the coin flips are  $Y_1, Y_2, Y_3, Y_4$  respectively. Based on the question description, for arbitrary  $i \in [1, 4]$  the event that a coin is “heads” is  $Y_i = 0$ . Therefore, we can call four flips as the random vector  $(Y_1, Y_2, Y_3, Y_4)$ . Understanding that selecting the standard coin is the event that  $X = 1$ , the sentence can be succinctly described as follows:

$$P(X = 1 | (Y_1, Y_2, Y_3, Y_4) = (0, 0, 0, 0))$$

### Part 2 Now, solve the expression.

For simplicity, we will let event  $E$  be the event that  $(Y_1, Y_2, Y_3, Y_4)$  is equal to  $(0, 0, 0, 0)$ , where again a zero indicates a heads.

We wish to find  $P(X = 1 | E)$ . We will implement Bayes’ Theorem.

$$P(X = 1 | E) = \frac{P(E | X = 1)P(X = 1)}{P(E)}$$

We can compute the probability of picking Coin #1 directly.

$$P(X = 1) = P(\text{unif}(\{0, 1, 2\}) = 1) = \frac{1}{3}$$

Assuming each  $Y_i$  is iid for  $i \in [1, 4]$ , we can compute  $P(E|X = 1)$ :

$$P(E|X = 1) = P((Y_1, Y_2, Y_3, Y_4) = (0, 0, 0, 0)|X = 1) = \prod_{i=1}^4 P(Y_i = 0|X = 1) = \left(\frac{1}{2^4}\right) = \frac{1}{16}$$

Now, we can compute  $P(E)$  via the Law of Total Probability. Trivially, we know  $P(E|X = 2) = 1$  and  $P(E|X = 0) = 0$ .

$$P(E) = \sum_{i=0}^2 P(E|X = i)P(X = i) = (0) \left(\frac{1}{3}\right) + \left(\frac{1}{16}\right) \left(\frac{1}{3}\right) + (1) \left(\frac{1}{3}\right) = \frac{1}{48} + \frac{1}{3} = \frac{17}{48}$$

Putting this all together, we have the following:

$$P(X = 1|E) = \frac{P(E|X = 1)P(X = 1)}{P(E)} = \frac{(1/3)(1/16)}{(17/48)} = \boxed{\frac{1}{17}}$$

### Question 3 : Non-Uniform Prior on Coin Types

We now modify the problem as follows: I stuffed the bag with 100 coins: 98 standard (fair) coins, 1 coin with only heads, and 1 coin with only tails. The rest is the same: pick one of the coins, flip it 4 times.

**Part 1** Write the joint distribution of this modified model. Use the notation as in Equation 1. Hint: use a **Categorical** distribution.

Now, we know that  $X$  is no longer Uniform. Using the notation for the categorical distribution used in class, we write that

$$X \sim \text{categorical}\left\{(0, 1, 2), \left(\frac{1}{100}, \frac{98}{100}, \frac{1}{100}\right)\right\}$$

Since the rules for flipping the coins remains the same, we still have that:

$$Y_i | X \sim \text{bern}(X/2)$$

Where  $X$  is now the categorical distribution.

**Part 2** Compute the probability that you picked one of the fair coins, given you see four heads.

**NOTE:** We will let the event of a heads be represented with a **1** rather than a **0**, to comply with the **dunif** function used in the succeeding questions.

We wish to find  $P(X = 1 | Y_{1:4} = \vec{\mathbf{1}})$ . We will implement Bayes' Theorem.

$$P(X = 1 | Y_{1:4} = \vec{\mathbf{1}}) = \frac{P(Y_{1:4} = \vec{\mathbf{1}}|X = 1)P(X = 1)}{P(Y_{1:4} = \vec{\mathbf{1}})}$$

We note from the Categorical Distribution defined above, that  $P(X = 1) = \frac{98}{100}$ . Similarly, by the same logic applied in **Question 2**, we know that:

$$P(Y_{1:4} = \vec{\mathbf{1}}|X = 1) = \left(\frac{1}{2}\right)^4 = \frac{1}{16}$$

Now, we can find  $P(Y_{1:4} = \vec{1})$  via the Law of Total Probability.

$$\begin{aligned}
 P(Y_{1:4} = \vec{1}) &= \sum_{i=0}^2 P(X = i)P(Y_{1:4} = \vec{1} \mid X = i) \\
 P(Y_{1:4} = \vec{1}) &= \frac{1}{100}(0) + \frac{98}{100}\left(\frac{1}{16}\right) + \frac{1}{100}(1) \\
 P(Y_{1:4} = \vec{1}) &= \frac{1}{100} + \frac{98}{1600} = \frac{114}{1600}
 \end{aligned}$$

Now, we can bring this all together to solve for  $P(X = 1 \mid Y_{1:4} = \vec{1})$

$$P(X = 1 \mid Y_{1:4} = \vec{1}) = \frac{P(Y_{1:4} = \vec{1} \mid X = 1)P(X = 1)}{P(Y_{1:4} = \vec{1})} = \frac{\left(\frac{1}{16}\right)\left(\frac{98}{100}\right)}{\left(\frac{114}{1600}\right)} = \boxed{\frac{49}{57}}$$

Where  $\frac{49}{57}$  is approximately 0.859649... as a decimal.

## Question 4

We now generalize to having  $K + 1$  types of coins such that:

- coin type  $k \in \{0, 1, \dots, K\}$  has bias  $k/K$
- the fraction of coins in the bag of type  $k$  is  $\rho_k$ .

We consider the same observation as before: “you observe 4 heads”. We want to find the conditional probability  $\pi_k$ , for all  $k$  that we picked coin type  $k \in \{0, 1, \dots, K\}$  from the bag given the observation.

**Part 1** Write an R function called `posterior_given_four_heads` taking as input a vector  $\rho = (\rho_0, \rho_1, \dots, \rho_K)$  and returning  $\pi = (\pi_0, \pi_1, \dots, \pi_K)$ .

```

posterior_given_four_heads <- function(rho){
  kval = 0:(length(rho)-1)
  K = length(kval)
  # kval = ( 0:(K-1) )
  biases = kval / (K-1)
  # in this case, the denominator is from LOTP again
  p_E = sum( biases^4 * rho )
  # then numerator is P(X)P(E | X)
  prior_likeli = sapply(1:K, function(k){ biases[k]^4 * rho[k] } )
  # then return the ratio
  return(prior_likeli / p_E)
}

# this returns q2
posterior_given_four_heads(rho = c(1/3, 1/3, 1/3))

```

```
## [1] 0.00000000 0.05882353 0.94117647
```


**Part 2** Test your code by making sure you can recover the answer in Q. 3 as a special case. Report what values of  $K$  and  $\rho$  you used.

**Part 3 (INCOMPLETE)** Show the output for  $\rho \propto (1, 2, 3, \dots, 10)$ . Here  $\propto$  means “proportional to”; try to infer what it means in this context.

```
posterior_given_four_heads((1:10))
```

```
## [1] 0.000000e+00 1.468882e-05 3.525316e-04 2.379588e-03 9.400843e-03  
## [6] 2.754153e-02 6.662848e-02 1.410714e-01 2.707443e-01 4.818667e-01
```

```
plot(posterior_given_four_heads((1:10)))
```



What does “proportional to” mean in this context?

```
# note that the whole 1 to 10 thing is non-normalized  
# so we know the proportionality to 1:10 is the same as  
# actually passing the following probabilities  
1:10/(sum(1:10))
```

```
## [1] 0.01818182 0.03636364 0.05454545 0.07272727 0.09090909 0.10909091  
## [7] 0.12727273 0.14545455 0.16363636 0.18181818
```

```
# also code output is the same  
posterior_given_four_heads( 1:10/(sum(1:10)))
```

```
## [1] 0.000000e+00 1.468882e-05 3.525316e-04 2.379588e-03 9.400843e-03  
## [6] 2.754153e-02 6.662848e-02 1.410714e-01 2.707443e-01 4.818667e-01
```

## Q.5: Generalizing Observations

We now generalize Q. 4 as follows: instead of observing 4 “heads” out of 4 observations, say we observe `n_heads` out of `n_observations`, where `n_heads` and `n_observations` will be additional arguments passed into a new R function.

**Part 1 (INCOMPLETE)** Write the joint distribution of this modified model. Use the  $\sim$  notation as in Equation 1. Hint: use a Binomial distribution.

**Part 2** Write an R function called `posterior` taking three input arguments in the following order: a vector  $\rho$  as in Q. 4, as well as two integers, `n_heads` and `n_observations`.

```
posterior <- function(rho, n_heads, n_observations){  
  # this part is the same  
  kvals = 0:(length(rho)-1)  
  K = length(kvals)  
  biases = kvals / (K-1)  
  # we need the overall probability of nheads  
  p_E = sum(dbinom(x = n_heads, size = n_observations, p = biases)*rho)  
  # then, across the k coins that can be picked  
  # then numerator is still P(X)P(E | X)  
  prior_likeli = sapply(1:K, function(k){  
    dbinom(x = n_heads, size = n_observations, p = biases[k]) * rho[k] } )  
  # then return the ratio  
  return(prior_likeli / p_E)  
}
```

```
posterior(c(1/3, 1/3, 1/3), 4, 4)
```

**Part 3 (MAYBE INCOMPLETE)**

```
## [1] 0.00000000 0.05882353 0.94117647
```

Test your code by making sure you can recover the answer in Q. 3 as a special case.

**Part 4 (MAYBE INCOMPLETE)** Show the output for  $\rho \propto (1, 2, 3, \dots, 10)$  and `n_heads = 2` and `n_observations = 10`.

```
posterior(1:10, 2, 10)
```

```
## [1] 0.000000e+00 1.628596e-01 3.357600e-01 2.934782e-01 1.516748e-01  
## [6] 4.771276e-02 8.024795e-03 4.870678e-04 2.795670e-06 0.000000e+00
```

```
plot(posterior(1:10, 2, 10))
```

