

STAT 447 Assignment 9

MCMC Hacking

Caden Hewlett

2024-04-05

Data Import

We import the primary data set inspired by Davidson-Pilon, (2013) below:

```
# main data vector
sms_data = c(
  13,24,8,24,7,35,14,11,15,11,22,22,11,57,11,19,29,6,19,
  12,22,12,18,72,32,9,7,13,19,23,27,20,6,17,13,10,14,6,
  16,15,7,2,15,15,19,70,49,7,53,22,21,31,19,11,18,20,12,
  35,17,23,17,4,2,31,30,13,27,0,39,37,5,14,13,22)
# data frame
df = data.frame(
  num_texts = sms_data,
  day = 1:length(sms_data)
)
```

Bayesian Model

We denote C to be the change point, selected uniformly from days $d \in \{1, 2, \dots, N\}$ where N is the number of observations. Then, there is a likelihood for days less than change point C and a likelihood for days above the change point.

We can denote the model as follows:

$$\begin{aligned}\lambda_1 &\sim \exp(1/100) \\ \lambda_2 &\sim \exp(1/100) \\ C &\sim \text{unif}(\{1, 2, \dots, N\}) \\ Y_d \mid C, \{\lambda_1, \lambda_2\} &\sim \text{pois}(\mathbb{1}[d < C]\lambda_1 + \mathbb{1}[d \geq C]\lambda_2)\end{aligned}$$

We will also refer to the $\{\lambda_1, \lambda_2\}$ pair as $\vec{\lambda}$.

We provide an implementation of the joint distribution of this model below.

```
# inputs are lambdas, C and y
log_joint = function(rates, change_point, y) {

  # Return log(0.0) if parameters are outside of the support
  if (rates[[1]] < 0 | rates[[2]] < 0 | change_point < 1 | change_point > length(y))
    return(-Inf)

  log_prior =
    dexp(rates[[1]], 1/100, log = TRUE) +
    dexp(rates[[2]], 1/100, log = TRUE)
```

```

log_likelihood = 0.0
for (i in 1:length(y)) {
  rate = if (i < change_point) rates[[1]] else rates[[2]]
  log_likelihood = log_likelihood + dpois(y[[i]], rate, log = TRUE)
}

return(log_prior + log_likelihood)
}

```

Question 1: A Custom MCMC Sampler

NOTE We will briefly justify the reasoning for irreducibility of each kernel (and thus the combination of kernels) in Part 1 and leave invariance for the next section. Further, when we define these kernels, we will be doing so in terms of likelihoods γ . However, for more precision with respect to the true implementation, we will prove π -invariance with respect to the log joint model defined above.

Part 1: Algorithm

We will begin by defining two separate kernels K_1 and K_2 for the rate parameters $\{\lambda_1, \lambda_2\}$ and the change-point parameter C , respectively. We then unify these kernels by defining a kernel mixture for a selection probability ρ .

Let C^* and C be the proposed and current cutoff times, respectively. Similarly, we let $\vec{\lambda}^*$ and $\vec{\lambda}$ be the proposed and current rates, respectively. We use $\vec{\lambda} = \{\lambda_1, \lambda_2\}$ and $\vec{\lambda}^* = \{\lambda_1^*, \lambda_2^*\}$ interchangeably. Further, in each case, we define the proposal function using the following notation:

$$q(\{\vec{\lambda}^*, C^*\} \mid \{\vec{\lambda}, C\}) \equiv q(\{\lambda_1^*, \lambda_2^*, C^*\} \mid \{\lambda_1, \lambda_2, C\})$$

While slightly obtuse, this is intended to reflect the fact that the `rates` input is a vector, and the likelihood defined in the original model relies upon all three parameters. Further, this style allows us to maintain the same general form for q regardless of the kernel, which we can use to expedite the proof in Question 2.

We begin by discussing K_1 . This kernel will only modify the $\vec{\lambda}$ parameter. Since each Poisson rate parameter is a real number, we will use a standard normal proposal.

Notably, rather than having a dimension of 2, we have each rate operate separately on their own means to be updated over time. The objective is to have more independent exploration for pre-and-post change point rates. Notably, our proposal vector is $\{\vec{\lambda}^*, C\}$, highlighting that in this kernel we do not select new values for the change point.

$$\begin{aligned}
q_1(\{\vec{\lambda}^*, C\} \mid \{\vec{\lambda}, C\}) &= \left\{ \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{1}{2}(\lambda_i^* - \lambda_i)^2\right) : i \in \{1, 2\} \right\} \\
\alpha_1(\vec{\lambda}^*, \vec{\lambda}, C) &= \min\left\{1, r(\vec{\lambda}^*, \vec{\lambda}, C)\right\}, \text{ where } r(\vec{\lambda}^*, \vec{\lambda}, C) = \frac{\gamma(\vec{\lambda}^*, C)}{\gamma(\vec{\lambda}, C)} \\
K_1(\{\vec{\lambda}^*, C\} \mid \{\vec{\lambda}, C\}) &= q_1(\{\vec{\lambda}^*, C\} \mid \{\vec{\lambda}, C\}) \cdot \alpha_1(\vec{\lambda}^*, \vec{\lambda}, C)
\end{aligned}$$

In a simplified version, we can write the equation below to simplify K_1 :

$$K_1(\vec{\lambda}^* \mid \vec{\lambda}) = q(\vec{\lambda}^* \mid \vec{\lambda}) \alpha(\vec{\lambda}^*, \vec{\lambda})$$

Further, we note that the standard normal proposal will allow the algorithm to explore all candidate $\vec{\lambda}$ pairs in $(\mathbb{R}^+)^2$. Let's talk briefly about why by highlighting an arbitrary proposed $\lambda_i^* \in \vec{\lambda}^*$ in terms of Poisson probability mass $\gamma(\vec{\lambda}^*, C)$ for well-defined $C \in [1, N] \subseteq \mathbb{N}$.

Firstly, we note that Poisson rate parameters λ are strictly positive, while the standard normal surrounding λ_i can propose $\lambda_i^* \in \mathbb{R}$. Notably, however, a property of the Poisson probability density (specifically for the `log_joint` function) is that:

$$p_{\text{pois}}(k; \lambda) = \mathbb{P}(\text{pois}(\lambda) = k) = \begin{cases} \frac{\lambda^k e^{-\lambda}}{k!}, & k \geq 0, \lambda \geq 0 \\ 0, & \text{otherwise} \end{cases}$$

Notably, in the case that either λ_1^* or λ_2^* are negative, the function will return zero. In this case the Metropolis-Hastings ratio will return zero, and the proposed $\vec{\lambda}^*$ will be rejected. Hence, we know that in K_1 that not only will solely valid $\vec{\lambda}$ be explored, but further assuming that C is well-defined all candidate lambda pairs have nonzero probability mass; hence, all possible $\vec{\lambda}^* \in (\mathbb{R}^+)^2$ can be explored from any given $\vec{\lambda}$ - thus this component of the kernel is irreducible. In the next section, we will verify the invariance of the overall kernel.

Now, we consider K_C . Ideally, we would like to define a symmetric proposal q_C ; however, we cannot use a Normal Distribution as before so we must find a symmetric discrete distribution. As per the recommendation, this distribution should have variance greater than 3 to avoid slow mixing.

The simplest possible proposal in this situation would be the discrete $\text{unif}(\{a = 1, b = N\})$ distribution, i.e.:

$$q(\{\vec{\lambda}, C^*\} | \{\vec{\lambda}, C\}) = \frac{1}{b - a + 1} = \frac{1}{N - 1 + 1} = \frac{1}{N}$$

Where, in this case, $N = 74$. This choice ensures that every possible change point C^* can be proposed with nonzero probability given any current C , preserving irreducibility.

We can verify by computing the variance of this distribution (to check it will mix nicely):

$$\text{var}(\text{unif}\{a, b\} | a = 1, b = N) = \frac{(b - a + 1)^2 - 1}{12} = \frac{(N - 1 + 1)^2 - 1}{12} = \frac{N^2 - 1}{12} \approx 456 \text{ for } N = 74$$

Which is certainly a large enough variance to have a breadth of proposal options.

Then, we define the kernel K_2 in a similar fashion to before.

$$\begin{aligned} q_2(\{\vec{\lambda}, C^*\} | \{\vec{\lambda}, C\}) &= \frac{1}{N} \\ \alpha_2(\vec{\lambda}, C, C^*) &= \min\{1, r(\vec{\lambda}, C, C^*)\}, \text{ where } r(\vec{\lambda}, C, C^*) = \frac{\gamma(\lambda, C^*)}{\gamma(\lambda, C)} \\ K_2(\{\vec{\lambda}, C^*\} | \{\vec{\lambda}, C\}) &= q_2(\{\vec{\lambda}, C^*\} | \{\vec{\lambda}, C\}) \cdot \alpha_2(\vec{\lambda}, C, C^*) \end{aligned}$$

Finally, with each kernel defined above we can define the unified kernel K as follows. This is what we will be implementing in Part 3.

We use the definitions of K_1 and K_2 detailed above.

Let $\rho \sim \text{bern}(1/2)$

$$K(\{\vec{\lambda}^*, C^*\} | \{\vec{\lambda}, C\}) = \mathbb{1}[\rho = 1]K_1(\{\vec{\lambda}^*, C\} | \{\vec{\lambda}, C\}) + (1 - \mathbb{1}[\rho = 1])K_2(\{\vec{\lambda}, C^*\} | \{\vec{\lambda}, C\})$$

It should be noted that the likelihoods discussed as part of K_1 and K_2 will be replaced by the log-joint in the software implementation, which slightly changes the structure of the acceptance criteria. In the next part, we discuss how this slight modification still preserves π -invariance.

Part 2: π -Invariance

Prove that the MCMC algorithm you defined in part 1 is π -invariant.

Since detailed balance implies global balance, it is sufficient to show that the DBEs hold to show invariance. We know that if each K_i in the kernel mixture is π -invariant that their mixture is π -invariant. Hence, we will separately prove π -invariance for K_λ and K_C hold under detailed balance to show that K is invariant.

Part 3: Implementation

Implement the MCMC algorithm you describe mathematically in R.

```
N = nrow(df)
# we build such that dim = 1
mcmc = function(rates, CP, y, n_iterations, debug = FALSE) {
  change_point_trace = rep(-1, n_iterations)
  # initial point
  current_CP = CP
  current_RT = rates
  current = list(rates, CP)
  # iteration station
  for (i in 1:n_iterations) {
    if (debug) {print(unlist(rep("#", times = 10)))}
    if (debug) {print(paste("Iteration:", i))}
    # bernoulli trial
    kernel_choice = ifelse(runif(1) < 0.5, 1, 2)
    if (debug) {print(paste("Chose Kernel", kernel_choice, "..."))}
    # Kernel for Change Point
    if (kernel_choice == 1){
      # Discrete Uniform Proposal
      prop = rdunif(1, 1, N)
      if (debug) {print(paste("Proposed", prop, "..."))}
      # MH ratio
      ratio = (log_joint(current[[1]], prop, y) -
               log_joint(current[[1]], current[[2]], y))
      # Bernoulli Trial
      if (log(runif(1)) < ratio) {
        # accept
        current[[2]] = prop
        if (debug) {print("Accepted!")}
      } else {
        # reject (redundant but nice)
        current[[2]] = current[[2]]
        if (debug) {print("Rejected!")}
      }
    }
    # Kernel for Lambdas
    else {
      # normal at current point
      ell_1 = rnorm(1, mean = current[[1]][1])
      ell_2 = rnorm(1, mean = current[[1]][2])
      # then the proposal is the vector
      prop = c(ell_1, ell_2)
      if (debug) {print(paste("Proposed", unlist(round(prop,2)), "..."))}
      # MH Ratio
      ratio = (log_joint(prop, current[[2]], y) -
               log_joint(current[[1]], current[[2]], y))
      if (log(runif(1)) < ratio) {
        # accept
        current[[1]] = prop
        if (debug) {print("Accepted!")}
      } else {
        # reject (redundant but nice)
      }
    }
  }
}
```

```

        current[[1]] = current[[1]]
        if (debug) {print("Rejected!")}
    }
    # update trace
    change_point_trace[i] = current[[2]]
    if(debug) {Sys.sleep(3)}
}
}
return(
  list(
    change_point_trace = change_point_trace,
    last_iteration_rates = current[[1]]
  )
)
}

```

```

## TESTING
set.seed(447)
# true change point at around 25
simulated_yvals = c(round(runif(24, 10, 15)), round(runif(60, 30, 50)))
# run MCMC
test = mcmc(c(0.1, 0.2), 34, simulated_yvals, 550, debug = FALSE)
# plot
plot(main = "Simulated Data Trace Plot",
     test$change_point_trace[test$change_point_trace > 0],
     type = 'l', ylab = "Proposal",
     ylim = c(1, 74))
abline(h = 25, col = rgb(1, 0, 0, 0.5), lwd = 1)

```

Simulated Data Trace Plot

